



# LAND OF THE CURIOUS





# TABLE OF CONTENTS

- » NoSQL databases
- » Benefits and drawbacks
- » Relational vs NoSQL
- » Column
- » Graph
- » Key-value
- » Document

 CT60A4304 - BASICS OF DATABASE SYSTEMS

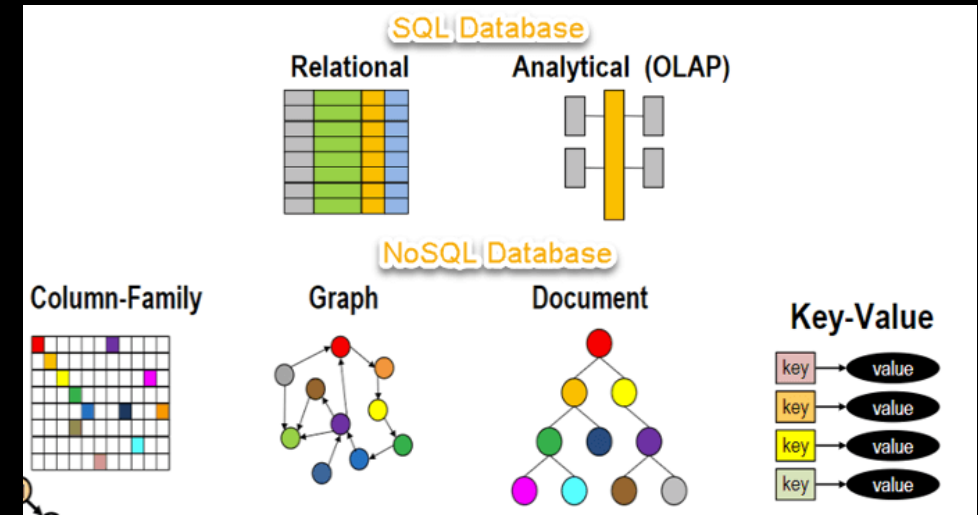
# NOSQL DATABASES

Lecture

Jiri Musto, D.Sc.

# NOSQL DATABASES

- » Stands for “Not Only SQL”
- » There are no fixed schemas
- » Avoids joins
- » Used for big data and real-time web apps
- » Can be structured, semi-structured or unstructured data
- » Developed in the 2000s







# BENEFITS AND DRAWBACKS

## » Benefits

- » Flexible
- » Scale-Out (distributed system)
- » Fast-paced development (no need for detailed modelling)
- » Semi-structured data

## » Drawbacks

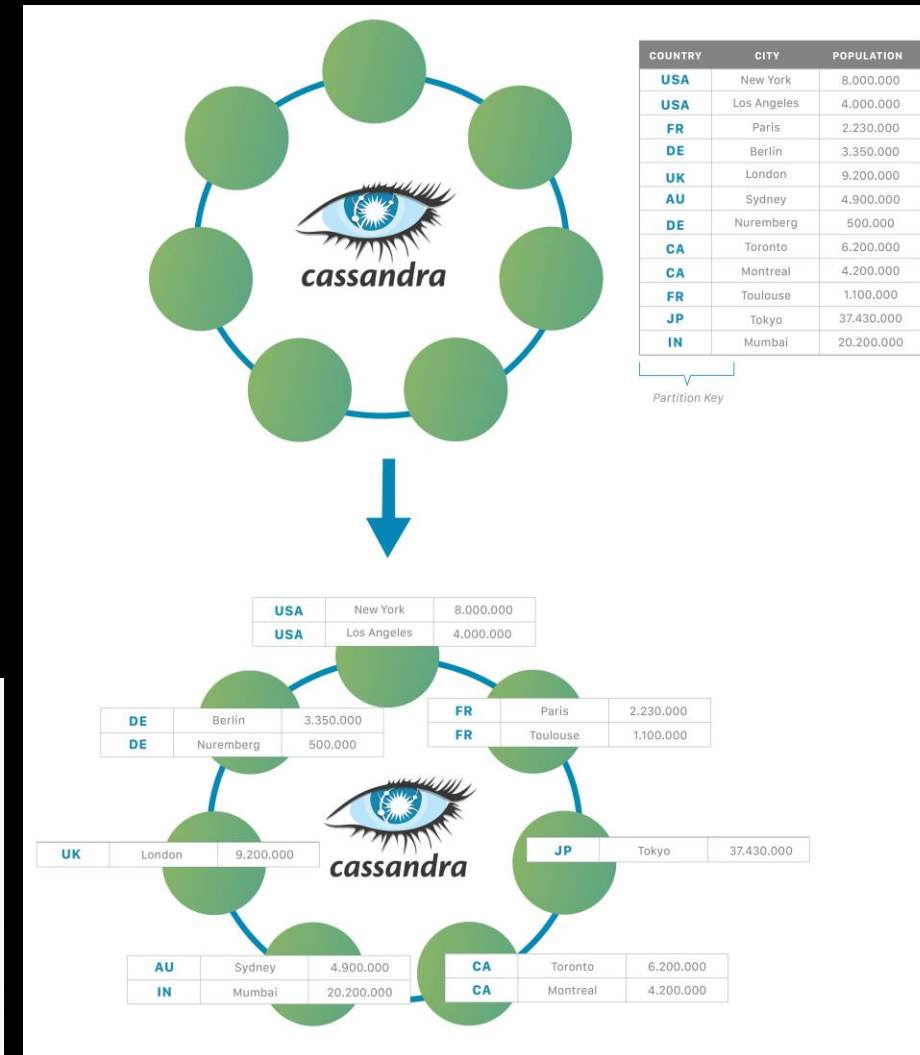
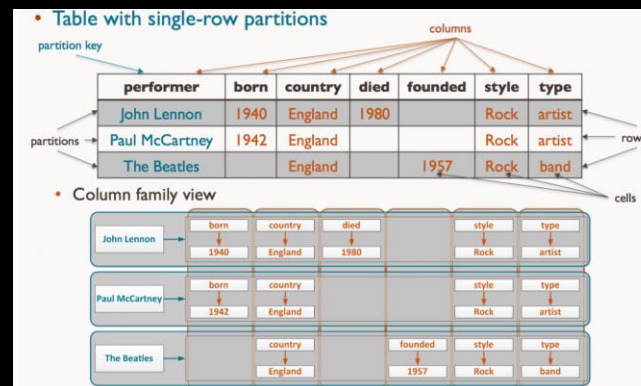
- » Relational database is better for scaling up (one system, more RAM, CPU, HDD)
- » Inconsistencies (not good for transaction oriented systems like bank)
- » Less secure
- » Lack of standardization

# RELATIONAL VS. NOSQL

Relational	NoSQL
Predefined schema	Schemaless
Transactions, consistent	Eventual consistency
Scale-up	Scale-out
Standard SQL language	No real standardization
Easy record retrieval	Easier to store large amounts of data
Structured data	Semi-structured data (or even unstructured)
Good for complex queries	Good for hierarchical data
Normalized schema	De-normalized schema

# COLUMN DATABASE

- » Useful in analytical applications or recommendation engines
- » Columns are treated separately
- » Can access rows with a “row key”
- » Example: Apache Cassandra and Apache Hbase
- » Use cases: Facebook, Spotify
- » Main benefit is partitioning data and aggregation functions



# GRAPH DATABASE

- » Stores the relationships between data as nodes
- » No need to join and match data from separate tables
- » Useful in social networks, product recommendations, access management, and similar scenarios
- » Example: Neo4j
- » Use cases: Walmart, Cisco







# KEY-VALUE DATABASE

- » Similar to Python dictionary, associative arrays, maps, hash, or other data structures
- » Each data has a key that directs to it
- » Value can be of any type (even data structures)
- » Used in session management and caching in web applications (login systems), user profiles, blog comments, e-commerce
- » Example: Redis, Dynamo
- » Use cases: Twitter timeline, Pinterest, Quora

# DOCUMENT DATABASE

- » Uses JSON, XML or BSON documents for storage
- » No predefined structure, documents can be different
- » Documents can have sub-documents or reference other documents
- » Accessing documents and fields within document works similar to key-value pairing
- » Can have their own query language as well as a universal API system
- » Useful in content management systems, e-commerce and real-time analytics
- » Example: CouchDB, MongoDB
- » Use cases: Twitter, Google, ebay

Col1	Col2	Col3	Col4
Data	Data	Data	Data
Data	Data	Data	Data
Data	Data	Data	Data

**Document 1**

```
{
  "prop1": data,
  "prop2": data,
  "prop3": data,
  "prop4": data
}
```

**Document 2**

```
{
  "prop1": data,
  "prop2": data,
  "prop3": data,
  "prop4": data
}
```

**Document 3**

```
{
  "prop1": data,
  "prop2": data,
  "prop3": data,
  "prop4": data
}
```

# EXAMPLE OF DOCUMENT DATABASE

```
{
  "_id": {
    "$oid": "620b26403845b9c5c7042819"
  },
  "uid": 84001001,
  "country_iso2": "US",
  "country_iso3": "USA",
  "country_code": 840,
  "fips": 1001,
  "county": "Autauga",
  "state": "Alabama",
  "country": "US",
  "combined_name": "Autauga, Alabama, US",
  "population": 55869,
  "loc": {
    "type": "Point",
    "coordinates": [-86.6441, 32.5395]
  },
  "date": {
    "$date": "2020-01-22T00:00:00.000Z"
  },
  "confirmed": 0,
  "deaths": 0,
  "confirmed_daily": 0,
  "deaths_daily": 0
}
```

```
{
  "_id": {
    "$oid": "620b26403845b9c5c704281a"
  },
  "uid": 84001001,
  "country_iso2": "US",
  "country_iso3": "USA",
  "country_code": 840,
  "fips": 1001,
  "county": "Autauga",
  "state": "Alabama",
  "country": "US",
  "combined_name": "Autauga, Alabama, US",
  "population": 55869,
  "loc": {
    "type": "Point",
    "coordinates": [-86.6441, 32.5395]
  },
  "date": {
    "$date": "2020-01-22T00:00:00.000Z"
  },
  "confirmed": 0,
  "deaths": 0,
  "confirmed_daily": 0,
  "deaths_daily": 0
}
```

```
_id: ObjectId("5e7a6f9f1f27c900177b30ea")
userId: ObjectId("5e7a6f9f1f27c900177b30eb")
prevVal: Object
  initialize: "New report"
newVal: Object
  initialize: "Report initialized"
systemTimeStamp: "Tue Mar 24 2020 20:37:51 GMT+0000 (Coordinated Universal Time)"
__v: 0
```

```
_id: ObjectId("5e7a70a21f27c900177b30ef")
additionalInfo: "You can use mobile phones to submit as well."
photo: "http://res.cloudinary.com/dtcshxdss/raw/upload/v1585082531/ffneygasssx..."
objectivityOfPath: 0
accuracyOfPath: 0
pathPhysicalPropertiesId: ObjectId("5e7a70a21f27c900177b30f1")
pathEnvironmentalPropertiesId: ObjectId("5e7a70a21f27c900177b30f0")
__v: 0
```

```

1  [{
2    "_id": {
3      "$oid": "5ea088738834618b60b586d1"
4    },
5    "created_at": "Wed Apr 22 18:09:11 +0000 2020",
6    "id": 1253023058413129700,
7    "id_str": "1253023058413129730",
8    "text": "@BobHoldenNYC Its where the corona breeds",
9    "truncated": false,
10   "entities": {
11     "hashtags": [],
12     "symbols": [],
13     "user_mentions": [
14       {
15         "screen_name": "BobHoldenNYC",
16         "name": "Robert Holden",
17         "id": 948549529615392800,
18         "id_str": "948549529615392769",
19         "indices": [
20           0,
21           13
22         ]
23       }
24     ],
25     "urls": []
26   },
27   "metadata": {
28     "iso_language_code": "en",
29     "result_type": "recent"
30   },
31   "source": "<a href=\"http://twitter.com/download/android\" rel=\"nofollow\">Twitter for Android</a>",
32   "in_reply_to_status_id": 1253012337667309600,
33   "in_reply_to_status_id_str": "1253012337667309574",
34   "in_reply_to_user_id": 948549529615392800,
35   "in_reply_to_user_id_str": "948549529615392769",
36   "in_reply_to_screen_name": "BobHoldenNYC",
37   "user": {
38     "id": 1042728872,
39     "id_str": "1042728872",
40     "name": "Rexxx",
41     "screen_name": "Rexxx0001",
42     "location": "New York, USA",
43     "description": "Hard rock/metal/drummer/scoundrel /animal lover/Yankee fan/gamer. #TheLucasOrder",
44     "url": null,

```

```

45   "entities": {
46     "description": {
47       "urls": []
48     }
49   },
50   "protected": false,
51   "followers_count": 115,
52   "friends_count": 331,
53   "listed_count": 1,
54   "created_at": "Fri Dec 28 17:32:20 +0000 2012",
55   "favourites_count": 14182,
56   "utc_offset": null,
57   "time_zone": null,
58   "geo_enabled": false,
59   "verified": false,
60   "statuses_count": 3350,
61   "lang": null,
62   "contributors_enabled": false,
63   "is_translator": false,
64   "is_translation_enabled": false,
65   "profile_background_color": "C0DEED",
66   "profile_background_image_url": "http://abs.twimg.com/images/themes/theme1/bg.png",
67   "profile_background_image_url_https": "https://abs.twimg.com/images/themes/theme1/bg.png",
68   "profile_background_tile": false,
69   "profile_image_url": "http://pbs.twimg.com/profile_images/1219784454941691904/UjMT6N8y_normal.jpg",
70   "profile_image_url_https": "https://pbs.twimg.com/profile_images/1219784454941691904/UjMT6N8y_normal.jpg",
71   "profile_banner_url": "https://pbs.twimg.com/profile_banners/1042728872/1570210305",
72   "profile_link_color": "1DA1F2",
73   "profile_sidebar_border_color": "C0DEED",
74   "profile_sidebar_fill_color": "DDEEFF",
75   "profile_text_color": "333333",
76   "profile_use_background_image": true,
77   "has_extended_profile": false,
78   "default_profile": true,
79   "default_profile_image": false,
80   "following": null,
81   "follow_request_sent": null,
82   "notifications": null,
83   "translator_type": "none"
84 },
85 "geo": null,
86 "coordinates": null,
87 "place": null,

```

This is one tweet from Twitter database

```

88 "coordinates": null,
89 "place": null,
90 "contributors": null,
91 "is_quote_status": false,
92 "retweet_count": 0,
93 "favorite_count": 0,
94 "favorited": false,
95 "retweeted": false,
96 "lang": "en"
97 }, {

```



# USING MONGODB IN A PROGRAM: EASY EXAMPLE

- » Setup the database connection
  - » Port, address, username, client
- » Connect to database
- » Insert into database
- » Read from database
  - » Within the program
  - » Or with external DBMS
    - MongoDB compass
    - MongoDB shell

```

22 //Port Setup... =====
23 const PORT = process.env.PORT || 5000;
24 //Uses =====
25 app.use(cors());
26 app.use(bodyParser.json());
27 app.use(express.static(path.join(__dirname, "client", "build")));
28
29 var MongoClient = require('mongodb').MongoClient;
30 var url = "mongodb://localhost:27017/";
31
32 function saveToDB(myobj, collectionName) {
33   MongoClient.connect(url, function(err, db) {
34     if (err) throw err;
35     var dbo = db.db("moreTweets");
36     dbo.collection(collectionName).insertOne(myobj, function(err, res) {
37       if (err) throw err;
38       db.close();
39     });
40   });
41 }
42
43 function tweetsFromNewYork(question, numberOfTweets) {
44   console.log("Twittering!!!");
45   clientApp.get(' https://api.twitter.com/1.1/search/tweets.json',
46     {
47       q : question,
48       geocode: "40.663471,-73.137798,75mi",
49       result_type : "recent",
50       lang: "en",
51       count: numberOfTweets
52     },
53     function(error, tweets, response) {
54       if(error) throw error;
55       tweets.statuses.forEach(element => {
56         saveToDB(element, "newYorkTweets");
57       });
58       console.log("Tweets saved!");
59     });
60 }
61

```

# USING A SCHEMA/MODEL WITH MONGODB

- » Middleware / library called *mongoose*
- » Make objects from JSON collections
- » Can assign data types, default values, restrictions, etc
  - » Similar to integrity constraints in SQL
- » Can also assign functions to mongoose models
  - » Similar to object-oriented programming

```
2 // grab the mongoose module
3 var mongoose = require('mongoose');
4
5 // define our address model
6 // module.exports allows us to pass this to other modules
7 module.exports = mongoose.model('Address',
8 {
9   street : {
10     type : String,
11     default: '',
12     required: true
13   },
14   city : {
15     type : String,
16     default: ''
17   },
18   geolocationId : {
19     type : mongoose.ObjectId,
20     require: true
21   },
22   precisionOfAddress : {
23     type : String,
24     default: '0'
25   },
26   accuracyOfLocation : {
27     type : String,
28     default: '0'
29   },
30   completenessOfLocation : {
31     type : Number,
32     default: 0
33   }
34 }, 'Address');
```

```
1 const kittySchema = new mongoose.Schema({
2   name: String
3 });
4
5 const Kitten = mongoose.model('Kitten', kittySchema);
6
7 const silence = new Kitten({ name: 'Silence' });
8 console.log(silence.name); // 'Silence'
9
10 // NOTE: You must add to the schema before compiling it with mongoose.model()
11 kittySchema.methods.speak = function speak() {
12   const greeting = this.name
13     ? "Meow name is " + this.name
14     : "I don't have a name";
15   console.log(greeting);
16 };
17
18 const Kitten = mongoose.model('Kitten', kittySchema);
19
20 const fluffy = new Kitten({ name: 'fluffy' });
21 fluffy.speak(); // "Meow name is fluffy"
22
```

 CT60A4304 - BASICS OF DATABASE SYSTEMS

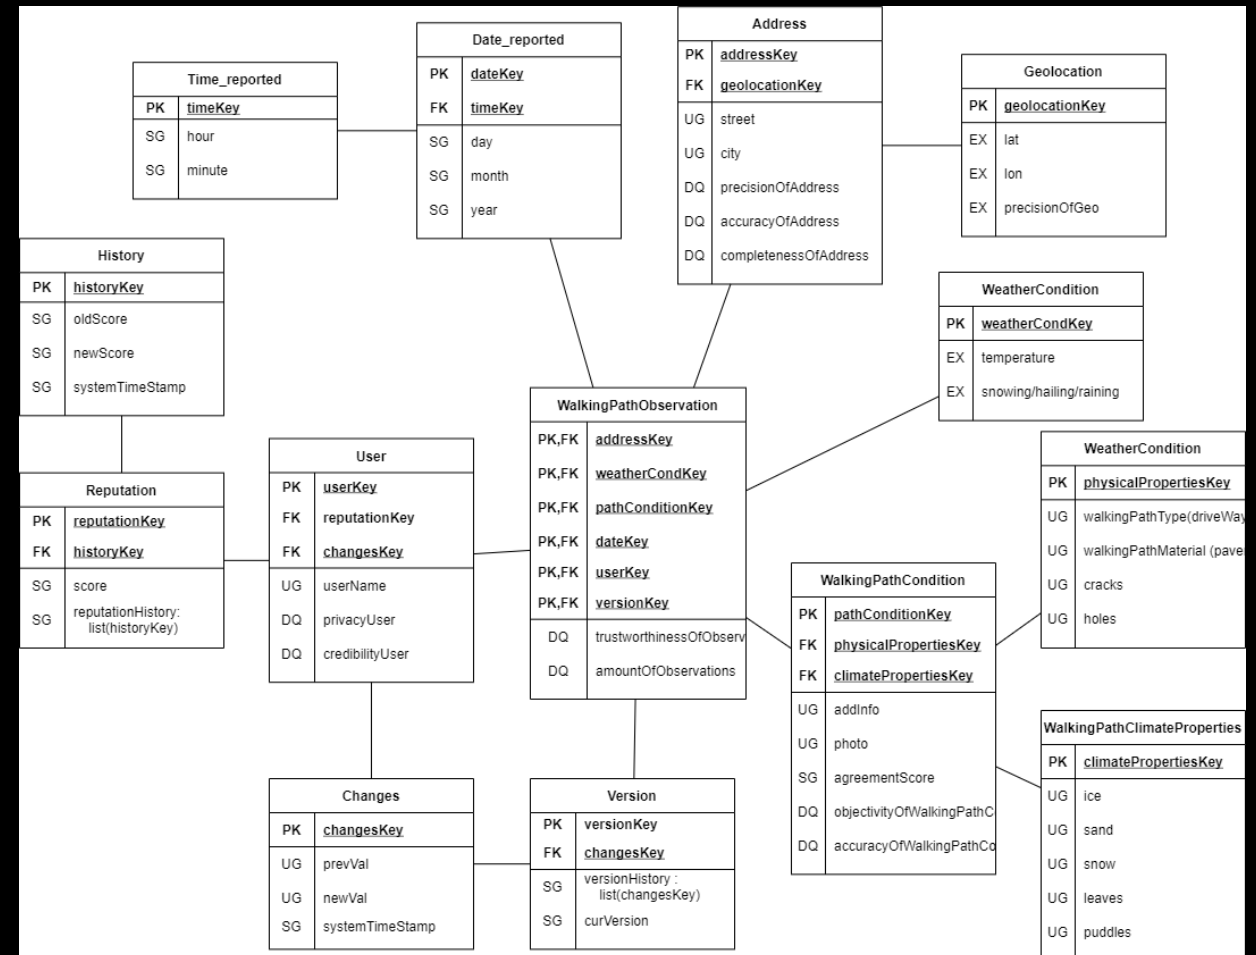
# MONGODB USE CASE

Lecture

Jiri Musto, D.Sc.

# DATABASE SCHEMA

- » Modeled as a snowflake schema
- » Acts like a relational database
- » Developed using MongoDB and mongoose
  - » Could have been done as a relational database
  - » Would have been easier as a relational database





# USING MONGODB: USING IT LIKE A RELATIONAL DATABASE

- » After setting connection
- » Setup mongoose models (optional)

```

13 let Address = require('../models/address.model');
14 let Condition = require('../models/condition.model');
15 let DateTime = require('../models/dateTimeSystem.model');
16 let GeoLocation = require('../models/geoLocationSystem.model');
17 let PathEnvironmentalProperties = require('../models/pathEnvironmentalProperties.model');
18 let PathPhysicalProperties = require('../models/pathPhysicalProperties.model');
19 let Report = require('../models/report.model');
20 let ReportVersion = require('../models/reportVersion.model');
21 let ReputationHistory = require('../models/reputationHistory.model');
22 let User = require('../models/user.model');
23 let UserChanges = require('../models/userChanges.model');
24 let UserReputation = require('../models/userReputation.model');
25 let WeatherSystem = require('../models/weatherSystem.model');
26 let apiTemperature = -200.0;
27 let testVariable = "not initialized";
28

```

# USING MONGODB: RETRIEVE ALL DATA

## » Search for the documents

```
/**Create the response array and template for one response */  
var responseArray = [];  
  
/**Get all reports and other data and their IDs */  
const reports = await Report.find(function(err,arr) {});  
const addresses = await Address.find(function (err,arr){});  
const dateTime = await DateTime.find(function (err,arr){});  
  
//const user = await User.findOne({_id:report.userId}, function(err,arr){});  
const weatherSystems = await WeatherSystem.find(function (err,arr){});  
const reportVersions = await ReportVersion.find(function (err,arr){});  
const geoLocations = await GeoLocation.find(function (err,arr){});  
const conditions = await Condition.find(function (err,arr){});  
const pathEnvironments = await PathEnvironmentalProperties.find(function (err,arr){});  
const pathPhysicals = await PathPhysicalProperties.find(function (err,arr){});
```

# USING MONGODB: CONTINUE

- » Collect all object ids to combine data into a tabular form in a later step

```
reports.forEach(function(item, index){
  var roadReport = item;
  oneReport = {
    reportId : "",
    dateTimeId : "",
    addressId : "",
    conditionId: "",
    weatherId: "",
    reportVersionId: "",
    date : "",
    time : "",
    street : "",
    city : "",
    delay: "",
    typeMaterial : "",
    physProperties : "",
    envProperties : "",
    geoPrecision: "",
    curVersion: "",
    reputationScore: "",
    trustworthinessOfReport: "",
    additonalInfo: "",
    temperature: "",
    season: "",
    raining: "",
    photo:""
  };
  /** After finding each report, build the response array and set the important IDs */
  oneReport.reportId = roadReport._id;
  oneReport.dateTimeId = roadReport.dateTimeId;
  oneReport.addressId = roadReport.addressId;
  oneReport.conditionId = roadReport.conditionId;
  oneReport.weatherId = roadReport.weatherId;
  oneReport.reportVersionId = roadReport.reportVersionId;
  oneReport.trustworthinessOfReport = roadReport.reputationScore;
  oneReport.reputationScore = roadReport.reputationScore;
  responseArray.push(oneReport);
});
```

# USING MONGODB: CONTINUE

- » Search for specific documents (objects) based on the object id

```
responseArray.forEach(function(report){
    /**
     * Based on the important IDs, find all relevant information
     * from the other arrays and build the response
     */
    dateTimes.forEach(function(date) {
        if(new String(report.dateTimeId).valueOf() == new String(date._id).valueOf()) {
            let d = date.day;
            let m = date.month;
            if (parseInt(date.day) < 10) {
                d = "0" + date.day;
            }
            if (parseInt(date.month) < 10) {
                m = "0" + date.month;
            }
            report.date = d + "." + m + "." + date.year;
            report.time = date.time;
            report.delay = date.delay;
        }
    });
    reportVersions.forEach(function(version) {
        if(new String(report.reportVersionId).valueOf() == new String(version._id).valueOf()) {
            report.curVersion = version.curVersion;
        }
    });
    weatherSystems.forEach(function(weather) {
        if(new String(report.weatherId).valueOf() == new String(weather._id).valueOf()) {
            report.temperature = parseFloat(weather.temperature);
            report.season = weather.season;
            report.raining = weather.raining;
        }
    });
    addresses.forEach(function(address) {
        if(new String(report.addressId).valueOf() == new String(address._id).valueOf()) {
            report.city = address.city;
            report.street = address.street;
            geoLocations.forEach(function(geoLocation) {
                if(new String(address.geoLocationId).valueOf() == new String(geoLocation._id).valueOf()) {
                    report.geoPrecision = geoLocation.geoPrecision;
                }
            });
        }
    });
});
```





## USING MONGODB: CONTINUE

»» Finally send the data to the client

```
res.json(responseArray);
```

# USING MONGODB: ADDING DATA

» Create new objects based on the models

```
let address          = new Address();

let condition        = new Condition();
let pathEnvironment  = new PathEnvironmentalProperties();
let pathPhysical     = new PathPhysicalProperties();

let geoLocation      = new GeoLocation();
let dateTime         = new DateTime();
let weather          = new WeatherSystem();

let report           = new Report();
let reportVersion    = new ReportVersion();
let userChanges      = new UserChanges();
```

# USING MONGODB: ADDING DATA

» Add the new data to the objects

```
dateTime.day = today.getDate();
dateTime.month = today.getMonth()+1; //0-11 months
dateTime.year = today.getFullYear();
let hours = today.getHours() + 2; //GMT +2 offset
let minutes = today.getMinutes();
let seconds = today.getSeconds();
if (today.getHours() < 10) {
    hours = "0" + (today.getHours() + 2);
}
if (today.getMinutes() < 10) {
    minutes = "0" + today.getMinutes();
}
if (today.getSeconds() < 10) {
    seconds = "0" + today.getSeconds();
}
dateTime.time = hours + ":" + minutes + ":" + seconds;

address.street = data.street;
address.city = data.city;
address.geoLocationId = geoLocation._id;
geoLocation.latitude = data.latitude;
geoLocation.longitude = data.longitude;
geoLocation.geoPrecision = data.geoPrecision;

weather.raining = "No rain";
weather.season = "Spring";
weather.temperature = apiTemperature;

condition.pathPhysicalPropertiesId = pathPhysical._id;
condition.pathEnvironmentalPropertiesId = pathEnvironment._id;
condition.additionalInfo = data.additionalInfo;
```

# USING MONGODB: ADDING DATA

- » Check if something exists (and act accordingly)

```
let user = null;
let existingUser = null;
try {
  existingUser = await User.findOne({userName: data.userName}, function (err, res){});
} catch (error) {
  existingUser = null;
}

if (existingUser == null) {
  user = new User();
} else {
  user = existingUser;
}
```



# USING MONGODB: ADDING DATA

» Save the objects to the database

```
address.save(function(err) {if (err) console.log(err); });  
geoLocation.save(function(err) {if (err) console.log(err); });  
dateTime.save(function(err) {if (err) console.log(err); });  
weather.save(function(err) { if (err) console.log(err); });  
report.save(function(err) { if (err) console.log(err); });  
condition.save(function(err) { if (err) console.log(err); });  
pathPhysical.save(function(err) { if (err) console.log(err); });  
pathEnvironment.save(function(err) { if (err) console.log(err); });  
userChanges.save(function(err) {if (err) console.log(err); });  
reportVersion.save(function(err) { if (err) console.log(err); });
```

# USING MONGODB: EDITING DATA

» Find the objects you want to edit

```
/**First find the one information based on the given ID */  
let report = await Report.findOne({_id:id}, function(err,arr){});  
let user = await User.findOne({userName:data.userName}, function(err,arr){});  
let address = await Address.findOne({_id:report.addressId}, function(err,arr){});  
let condition = await Condition.findOne({_id:report.conditionId}, function(err,arr){});  
let reportVersion = await ReportVersion.findOne({_id:report.reportVersionId}, function(err,arr){});  
  
let pathEnvironment = await PathEnvironmentalProperties.findOne({_id:condition.pathEnvironmentalPropertiesId}, function(err,arr){});  
let pathPhysical = await PathPhysicalProperties.findOne({_id:condition.pathPhysicalPropertiesId}, function(err,arr){});
```

# USING MONGODB: EDITING DATA

»» Update the given data

```
address.street = data.street;
address.city = data.city;
condition.additionalInfo = data.additionalInfo;

pathPhysical.pathType = data.pathType;
pathPhysical.pathMaterial = data.pathMaterial;
pathPhysical.cracks = data.cracks;
pathPhysical.holes = data.holes;
pathPhysical.debris = data.debris;

pathEnvironment.ice = data.ice;
pathEnvironment.snow = data.snow;
pathEnvironment.sand = data.sand;
pathEnvironment.puddles = data.puddles;
pathEnvironment.leaves = data.leaves;
pathEnvironment.pollen = data.pollen;
```



# USING MONGODB: EDITING DATA

» Save the new data

```
/**
 * All the save functions
 */
address.save(function(err) {if (err) res.send(err); });
report.save(function(err) { if (err) res.send(err); });
condition.save(function(err) { if (err) res.send(err); });
pathPhysical.save(function(err) { if (err) res.send(err); });
pathEnvironment.save(function(err) { if (err) res.send(err); });
reportVersion.save(function(err) { if (err) res.send(err); });
```



# FINAL NOTES

- » Using a NoSQL database in the same fashion as a relational database has its downsides
  - » Throwing away the benefits of a NoSQL and doing an inferior version of a relational database
- » The major benefit of a NoSQL database:
  1. Asynchronous operations
    - » User can continue using the interface while database retrieves or sends data
    - » In the case of failure, manual checks and rollbacks have to be done
  2. Modifying the models on during development



